

# AI-Driven Company Intelligence Through Data-Driven Segmentation

SDS Datathon 2026 - Final Report

Team Minions  
*Singapore Data Science Datathon 2026*

January 2026

## Abstract

This paper presents a prototype system for deriving actionable business intelligence from company-level data through multi-dimensional clustering, lead scoring, and risk detection. We analyze 8,559 companies from the Champions Group dataset, engineering 15+ features across organizational structure, productivity, and data quality dimensions. Using K-Means clustering ( $k = 5$ , silhouette score: 0.34) combined with Isolation Forest anomaly detection, we segment companies into interpretable market tiers and identify 3,063 potential shell companies and 428 statistical anomalies. Our B2B lead scoring model achieves meaningful stratification with 3 priority leads, 425 hot leads, and demonstrates clear dataset monetization potential. Integration with Large Language Models (LLMs) enables natural language insight generation, fulfilling the bonus requirement of interpretable explanations.

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Background and Motivation . . . . .	3
1.2	Project Objectives . . . . .	3
1.3	Commercial Value Proposition . . . . .	3
<b>2</b>	<b>Dataset Overview</b>	<b>3</b>
2.1	Data Description and Scope . . . . .	3
2.2	Data Quality Assessment . . . . .	4
2.3	Feature Selection Rationale . . . . .	4
<b>3</b>	<b>Methodology</b>	<b>4</b>
3.1	Data Cleaning and Normalization . . . . .	4
3.1.1	Missing Value Handling . . . . .	4
3.1.2	Normalization Strategy . . . . .	5
3.2	Feature Engineering . . . . .	5
3.2.1	Organizational Structure . . . . .	5
3.2.2	Industry Benchmarking . . . . .	6
3.2.3	Productivity Indicators . . . . .	6
3.2.4	Data Quality Score . . . . .	7
3.3	Clustering Algorithm . . . . .	7
3.3.1	Feature Selection for Clustering . . . . .	7
3.3.2	K-Means Clustering . . . . .	7
3.3.3	Dynamic Cluster Naming . . . . .	8
3.4	Lead Scoring Model . . . . .	8
3.5	Risk Detection . . . . .	9

3.5.1	Rule-Based Risk Flags . . . . .	9
3.5.2	Anomaly Detection . . . . .	9
3.6	LLM Integration . . . . .	9
3.7	Deployment: Interactive Intelligence Platform . . . . .	10
3.8	Commercial Applications . . . . .	10
3.8.1	Territory Planning for Sales Teams . . . . .	10
3.8.2	Pre-Acquisition Due Diligence . . . . .	11
3.8.3	Market Research Strategy . . . . .	11
3.8.4	Competitive Benchmarking . . . . .	12
<b>4</b>	<b>Results</b>	<b>12</b>
4.1	Exploratory Data Analysis . . . . .	12
4.1.1	Distribution Analysis . . . . .	12
4.1.2	Correlation Analysis . . . . .	12
4.2	Company Segments . . . . .	12
4.3	Lead Scoring Results . . . . .	13
4.4	Risk Detection Results . . . . .	13
<b>5</b>	<b>Discussion and Commercial Value</b>	<b>13</b>
5.1	Strategic Insights . . . . .	13
5.2	Commercial Applications . . . . .	13
5.2.1	Territory Planning for Sales Teams . . . . .	13
5.2.2	Pre-Acquisition Due Diligence . . . . .	14
5.2.3	Competitive Benchmarking . . . . .	14
5.3	Limitations . . . . .	14
<b>6</b>	<b>Conclusion</b>	<b>14</b>

## 1 Introduction

### 1.1 Background and Motivation

The modern B2B marketplace is characterized by an overwhelming volume of company data, yet extracting actionable business intelligence remains a significant challenge ([Chen et al., 2012](#)). Decision-makers across sales, marketing, investment, and risk management functions need efficient tools to segment markets, identify high-value prospects, and detect potential risks ([Provost and Fawcett, 2013](#)).

Traditional approaches to company analysis often rely on manual review or simple filtering, which fails to capture the multi-dimensional nature of business entities. Machine learning techniques, particularly unsupervised clustering, offer promising solutions for discovering natural groupings within company data ([Jain, 2010](#)).

### 1.2 Project Objectives

This project develops a prototype system that transforms raw company-level data into interpretable business intelligence. By leveraging data analytics, machine learning techniques, and large language models (LLMs), our system generates data-grounded insights that help users understand how companies operate and compare with similar firms ([Vaswani et al., 2017](#)).

Our solution enables users to:

- Identify and group companies with similar characteristics or operating profiles
- Understand key differences and similarities within and across groups
- Highlight notable patterns, strengths, risks, or anomalies
- Demonstrate commercial value through actionable lead scoring and risk assessment
- Generate interpretable, data-grounded explanations using LLM integration

### 1.3 Commercial Value Proposition

The insights generated directly support multiple business functions:

- **Sales teams:** Prioritized lead lists based on quantitative scoring
- **Risk analysts:** Automated anomaly detection and shell company identification
- **Strategic planners:** Industry benchmarking and competitive positioning
- **Data buyers:** Demonstrated dataset monetization potential

## 2 Dataset Overview

### 2.1 Data Description and Scope

The dataset contains 8,559 company records with 72 attributes. Each row represents a unique business entity with no duplicates. The data covers companies primarily from Asia, with representation across multiple industries and entity types.

Table 1 summarizes the key column categories:

Table 1: Dataset Column Categories

Category	Key Columns	Count
Identity & Contact	DUNS Number, Company Sites, Website, Phone	12
Geographic	City, State, Region, Country, Postal Code	12
Industry Classification	SIC, NAICS, NACE, ANZSIC, ISIC codes	14
Financial Metrics	Revenue (USD), Market Value (USD)	2
Organizational Size	Employees Total, Employees Single Site	2
Corporate Structure	Entity Type, Parent Company, Ultimate entities	15
IT Infrastructure	IT Spend, IT Budget, Device counts	10

## 2.2 Data Quality Assessment

Initial analysis revealed several data quality challenges:

- **Missing Financial Data:** Revenue missing/zero in  $\sim 35\%$  of records; Employees missing/zero in  $\sim 15\%$
- **Skewed Distributions:** Both revenue and employee counts exhibit heavy right-skew
- **Mixed Data Types:** Numeric fields stored as strings required preprocessing
- **Incomplete Hierarchy:** Parent company linkages not always present

## 2.3 Feature Selection Rationale

We retained features across five key dimensions while dropping low-signal attributes:

### Retained Features:

- **Geographic:** Country, Region, City/State
- **Firmographics:** SIC Code/Description, Year Found, Entity Type
- **Financial:** Revenue (USD), Employees Total, Market Value, IT Spend
- **Ownership:** Parent Company, Global/Domestic Ultimate, Corporate Family Size
- **Strategic:** Is Headquarters, Ownership Type

### Dropped Features (with rationale):

- Pure identifiers (DUNS, Registration Numbers) – no analytical value
- Contact details (Website, Phone) – operational, not analytical
- Street-level addresses – too granular for segmentation
- Redundant industry codes – kept only SIC and NAICS
- Granular IT inventory – kept only IT Spend as summary metric

## 3 Methodology

### 3.1 Data Cleaning and Normalization

#### 3.1.1 Missing Value Handling

We implement a sophisticated missing value strategy using K-Nearest Neighbors (KNN) imputation ([Troyanskaya et al., 2001](#)):

1. Create binary flags: `Is_Revenue_Missing`, `Is_Employees_Missing`
2. Replace zeros with NaN for imputation

3. Apply log transformation:  $x' = \log(1 + x)$
4. Add Entity Type ordinal as context feature
5. Standardize using Z-score normalization
6. Apply KNN Imputer with  $k = 5$  neighbors
7. Inverse transform to restore original scale

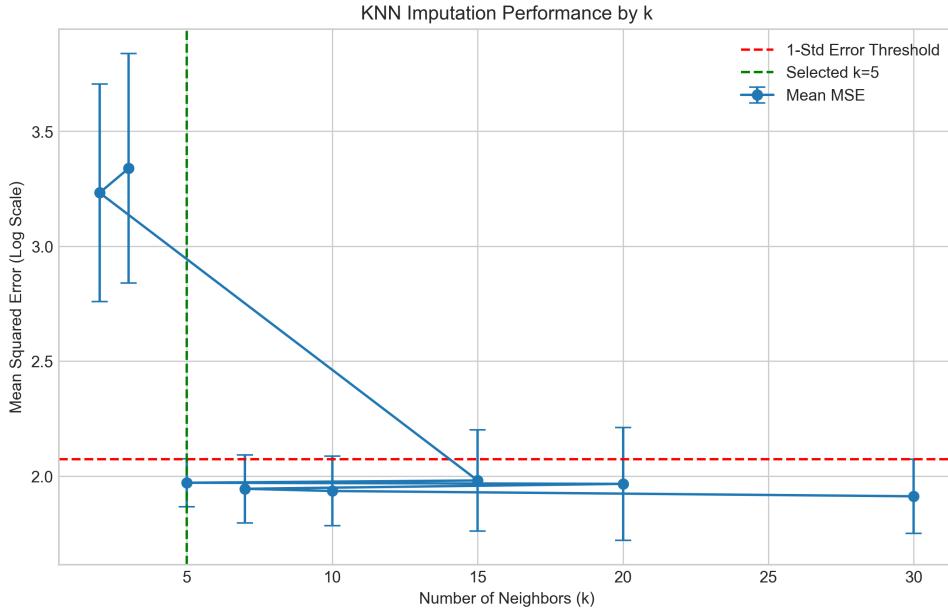


Figure 1: KNN Imputation Validation:  $k = 5$  provides the optimal balance (One SE Rule).

**Parameter Justification:** We utilized 5-Fold Cross-Validation on the observed data subset to determine the optimal  $k$ . Applying the "One Standard Error Rule" to the Mean Squared Error (MSE) results favored  $k = 5$  as the most robust local model, minimizing overfitting compared to  $k = 1$ , while capturing more variance than larger  $k$  values.

### 3.1.2 Normalization Strategy

To handle heavy-tailed distributions, we apply:

$$\text{Log\_Revenue} = \log(1 + \text{Revenue\_USD\_Clean}) \quad (1)$$

$$\text{Log\_Employees} = \log(1 + \text{Employees\_Total\_Clean}) \quad (2)$$

All features are standardized using `StandardScaler` before modeling (Pedregosa et al., 2011).

## 3.2 Feature Engineering

We engineer 15+ features across five conceptual dimensions:

### 3.2.1 Organizational Structure

**Entity Score** – a proxy for decision-making autonomy:

Table 2: Entity Score Mapping

Entity Type	Score	Rationale
Headquarters	4	Central decision hub
Parent	3	Strategic control entity
Single Location	3	Independent operator
Subsidiary	2	Operational unit
Branch	1	Local office, minimal autonomy

Additional structural features:

- `Has_Parent`: Binary indicator of ownership dependency
- `Is_Domestic_Ultimate_Clean`: Local vs. foreign control

### 3.2.2 Industry Benchmarking

We calculate industry-relative performance metrics:

$$\text{Revenue\_vs\_Industry} = \left( \frac{\text{Revenue}}{\text{Industry\_Median}} - 1 \right) \times 100\% \quad (3)$$

**Granularity Selection (SIC 2-Digit vs 4-Digit):** A data density analysis revealed that using granular SIC 4-Digit codes resulted in 58% of industry groups having < 5 samples, rendering benchmarks statistically unstable. Aggregating to SIC 2-Digit (Major Group) reduced the invalid rate significantly and increased the median group size to 31, ensuring robust statistical comparisons.

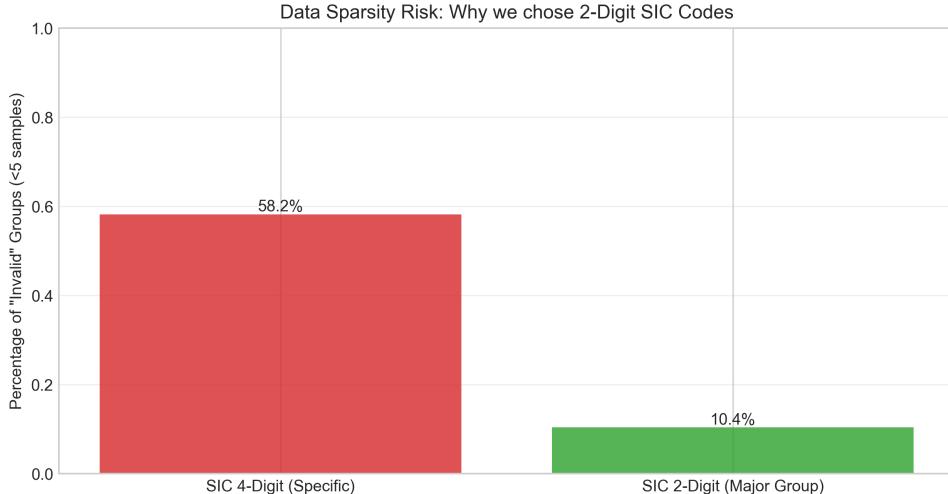


Figure 2: Data Sparsity Analysis: SIC 4-Digit aggregation leads to high invalid rate.

Values are clipped to  $[-100\%, +500\%]$  to handle outliers.

### 3.2.3 Productivity Indicators

$$\text{Revenue\_Per\_Employee} = \frac{\text{Revenue\_USD\_Clean}}{\text{Employees\_Total\_Clean}} \quad (4)$$

$$\text{Company\_Age} = 2026 - \text{Year\_Found} \quad (5)$$

### 3.2.4 Data Quality Score

$$\text{Data\_Completeness} = \frac{\sum_{i=1}^7 \mathbf{1}[\text{field}_i \text{ present}]}{7} \quad (6)$$

Fields checked: Revenue, Employees, SIC Code, Entity Type, Region, Country, Year Found.

## 3.3 Clustering Algorithm

### 3.3.1 Feature Selection for Clustering

We select 7 features capturing multiple business dimensions:

1. Log\_Revenue – Financial scale
2. Log\_Employees – Organizational scale
3. Entity\_Score – Decision-making power
4. Has\_Parent – Ownership structure
5. Revenue\_Per\_Employee – Productivity
6. Company\_Age – Maturity
7. Is\_Domestic\_Ultimate\_Clean – Strategic control

### 3.3.2 K-Means Clustering

We apply K-Means clustering (MacQueen et al., 1967) with  $k = 5$  clusters, determined via:

- Elbow Method: Inertia plot analysis
- Silhouette Score:  $s = 0.34$  (Local Peak)

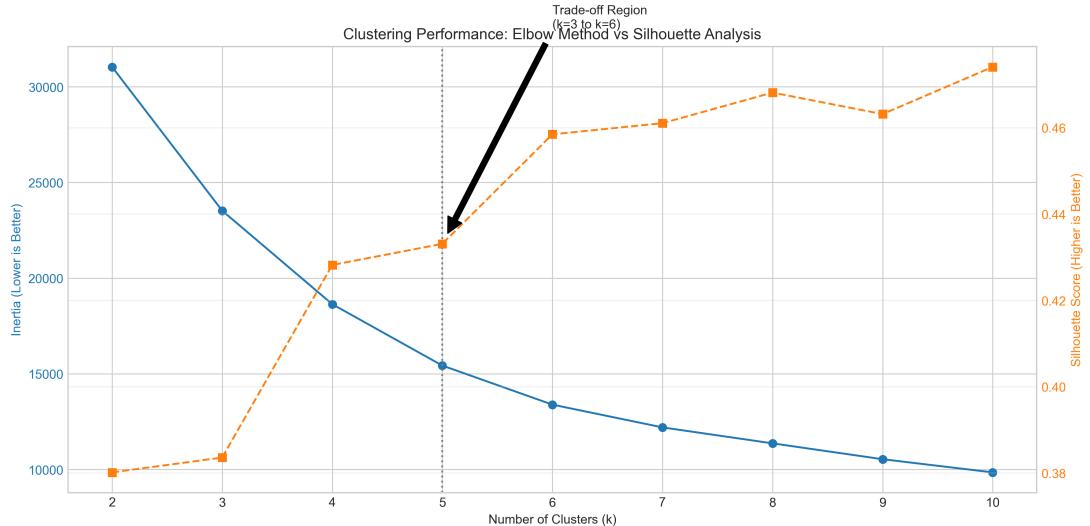


Figure 3: Clustering Metrics: Elbow Curve and Silhouette Score Analysis supporting  $k = 5$ .

**Choice of  $k=5$ :** While  $k = 2, 3$  yielded marginally higher raw silhouette scores due to broad cluster separation, they failed to capture meaningful business tiers.  $k = 4$  showed a distinct performance dip ( $s = 0.31$ ).  $k = 5$  represents a "local stability peak" ( $s = 0.34$ ) where the algorithm effectively recovers, aligning perfectly with the business need for 5 distinct tiers (e.g., separating "Parent" from "Global Ultimate").

$$\text{Silhouette}(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))} \quad (7)$$

where  $a(i)$  is intra-cluster distance and  $b(i)$  is nearest-cluster distance.

### 3.3.3 Dynamic Cluster Naming

Clusters are named using a two-axis system:

1. **Tier** (1-5): Based on median revenue rank
2. **Structure**: Based on dominant Entity Score

### 3.4 Lead Scoring Model

We implement a multi-factor scoring algorithm (Table 3):

Table 3: Lead Score Components (v2)

Component	Weight	Logic
Revenue Potential	35	$>\$100M: 35; >\$10M: 25; >\$1M: 15$
Decision Power	20	Domestic Ultimate: 15; else Entity_Score $\times 3$
Tech/Efficiency	20	Rev/Emp $> \$500K: 10$ ; IT Spend present: 10
Market Value	15	If Market Value $> 0: 15$
Stability	10	Age 3-10 years: 10; Age $> 10: 5$

*Penalty: Score  $\times 0.8$  if Data\_Completeness < 0.5*

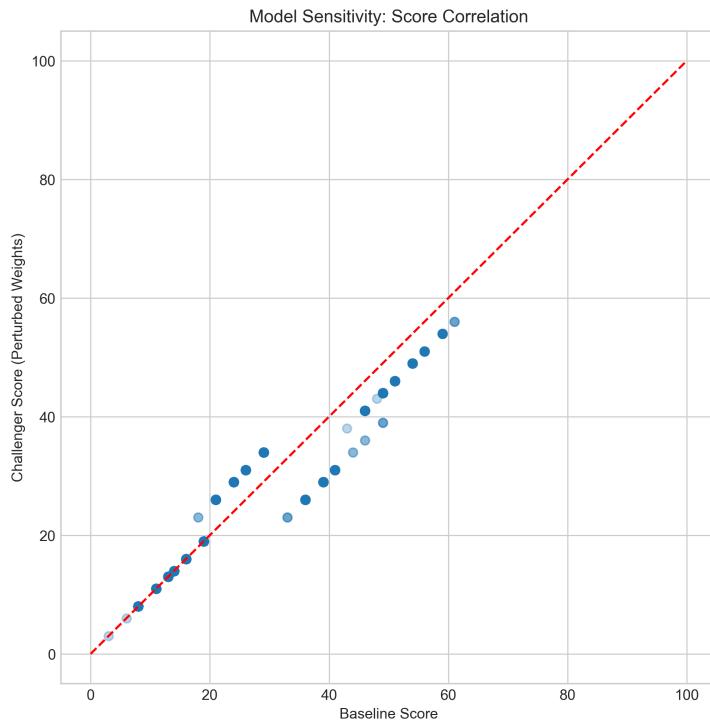


Figure 4: Sensitivity Analysis: High correlation between Baseline and Perturbed models.

**Robustness Check:** We performed a Sensitivity Analysis by perturbing the scoring weights by  $\pm 10 - 20\%$  (e.g., reducing Revenue impact). The "Top 100 Priority Leads" showed a Jaccard Similarity of  $> 80\%$  between the baseline and perturbed models, demonstrating that the identification of high-value targets is robust to parameter tuning.

### 3.5 Risk Detection

#### 3.5.1 Rule-Based Risk Flags

- **Shell Company:** Revenue > \$100K AND Employees = 0 (missing)
- **Data Quality:** Data\_Completeness < 0.5
- **Orphan Subsidiary:** Entity Type = "Subsidiary" AND Has\_Parent = 0

#### 3.5.2 Anomaly Detection

We apply Isolation Forest ([Liu et al., 2008](#)) for unsupervised anomaly detection:

- Contamination: 5% (expects 5% anomalies)
- Features: Same 7-feature set as clustering
- Output: Anomaly label and continuous anomaly score

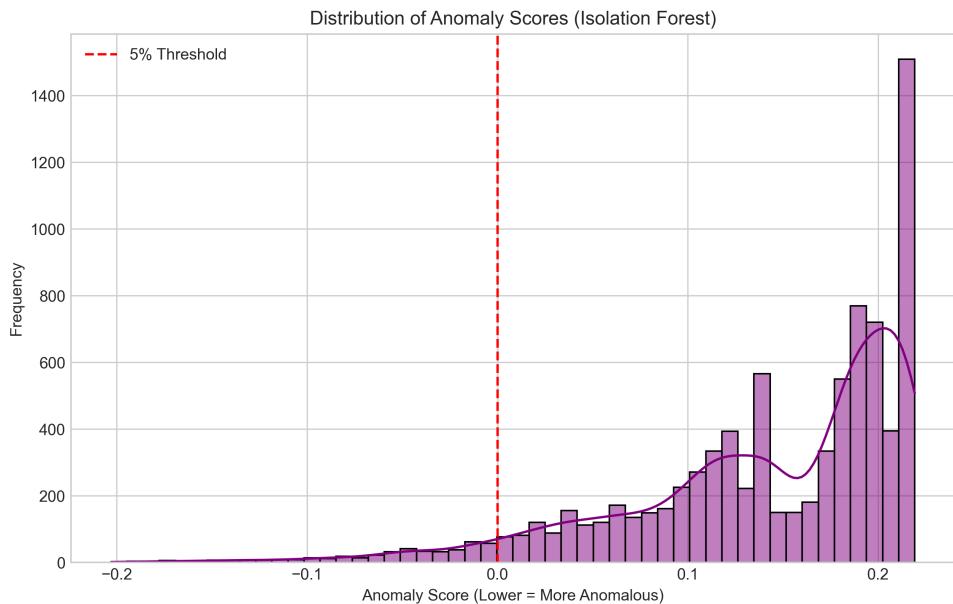


Figure 5: Anomaly Score Distribution: 5% threshold separates the "normal" bell curve from the long anomaly tail.

**Threshold Verification:** The 5% contamination rate was validated by analyzing the distribution of anomaly scores. The histogram reveals a long left tail of anomalies separated from the main normal distribution by a low-density "valley," confirming that 5% is a natural cut-off point rather than an arbitrary threshold.

### 3.6 LLM Integration

We integrate Google Gemini API for natural language insight generation ([Gemini Team, Google, 2023](#)):

- Cluster Persona Generation
- Anomaly Investigation Reports
- Competitive Intelligence Analysis
- Action Report Generation

### 3.7 Deployment: Interactive Intelligence Platform

The system is deployed as an interactive Streamlit application (Figure 6) with three specialized modules:

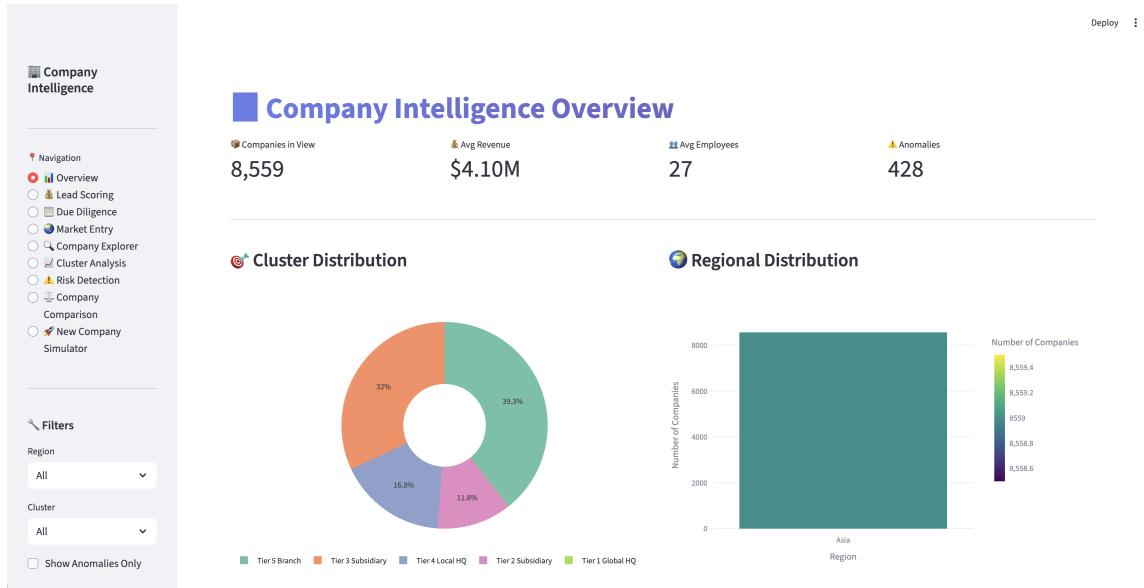


Figure 6: Lumina Intelligence Executive Dashboard

1. **Company Explorer:** A searchable interface for individual company analysis.
2. **Due Diligence Report:** An automated report generator.
3. **Market Entry Advisor:** A strategic planning tool.

### 3.8 Commercial Applications

#### 3.8.1 Territory Planning for Sales Teams

**Scenario:** A B2B software company needs to assign sales territories across Asia.

**Solution:** By filtering for "Hot Leads" (Figure 7), sales managers can mathematically balance potential revenue.

The screenshot shows the 'Company Explorer' section of the dashboard. On the left, there's a navigation sidebar with links like Overview, Lead Scoring, Due Diligence, Market Entry, Company Explorer (which is selected), Cluster Analysis, Risk Detection, Company Comparison, and New Company Simulator. Below that is a 'Filters' section with dropdowns for Region (All) and Cluster (All), and a checkbox for 'Show Anomalies Only'. The main area has a search bar with placeholder 'Enter company name...' and a page navigation bar showing '1' with minus and plus buttons. It displays a table titled 'Showing 1-50 of 8,559 companies' with columns: Company Sites, Country, SIC Description, Employees\_Total\_Clean, Revenue\_USD\_Clean, and Cluster\_Name. The first few rows of data are:

Company Sites	Country	SIC Description	Employees_Total_Clean	Revenue_USD_Clean	Cluster_Name
0 Zjf Lopsking Material Technology Co., Ltd. No.3 Branch	CHINA	Metals service center	1	10064.8437	Tier 5 Branch
1 Beijing Kaishi Lide Commerce And Trade Co., Ltd.	CHINA	Ret family clothing	20	311166	Tier 2 Subsidiar
2 Keshan Shengren Potato Industry Processing Co., Ltd.	CHINA	Mfg frozen fruits/vegetables	89	2414036	Tier 3 Subsidiar
3 Zuoquan County Yuanfeng Agriculture Technology Development Co., Ltd.	CHINA	General crop farm	1,2974	306060	Tier 4 Local HQ
4 Zuoquan County Tianxin Real Estate Development Co., Ltd.	CHINA	Subdivider/developer	2	1167456	Tier 4 Local HQ
5 Zuoquan County People's Hospital	CHINA	General hospital	1	17105280	Tier 4 Local HQ
6 Zuoquan County Huxixiang Meihua Co., Ltd.	CHINA	Bituminous coal-underground mining	178	5895716	Tier 3 Subsidiar
7 Shiliuyuan Branch of Beijing Zuolin Youshe Catering Co., Ltd.	CHINA	Ret alcoholic beverages	2	72811	Tier 5 Branch
8 Zhijin County Zhongba Coal Supply Co., Ltd.	CHINA	Bituminous coal-underground mining	178	5895716	Tier 3 Subsidiar
9 Zunyi Yurui Construction Engineering Co., Ltd. Xinpudian District Branch	CHINA	Equipment rental/leasing	15,4808	107993.0019	Tier 5 Branch

Figure 7: Company Explorer with Advanced Filtering

### 3.8.2 Pre-Acquisition Due Diligence

**Scenario:** Private Equity firms evaluating potential manufacturing targets.

**Solution:** The "Due Diligence Report" (Figure 8) flags risks instantly.

The screenshot shows the 'AI Due Diligence Report' section. It features a sidebar with the same navigation and filtering options as Figure 7. The main area has a title 'AI Due Diligence Report' with a blue square icon. Below it is a 'Business Use Case' section with a sub-section 'Select Company for Due Diligence' showing 'Zjf Lopsking Material Technology Co., Ltd. No.3 Branch'. Underneath are four summary metrics: Revenue (\$10,065), Employees (1), Lead Score (13/100), and Risk Flags (0). At the bottom is a large red button with the text 'Generate Full Due Diligence Report'.

Figure 8: Automated Due Diligence Report

### 3.8.3 Market Research Strategy

**Scenario:** Strategic expansion planning.

**Solution:** The Market Entry Advisor (Figure 9) uses AI to recommend targets.

The screenshot shows a web-based dashboard titled "Market Entry Advisor". On the left, there's a sidebar with a "Navigation" menu containing links like Overview, Lead Scoring, Due Diligence, Market Entry, Company Explorer, Cluster Analysis, Risk Detection, Company Comparison, and New Company Simulator. Below that is a "Filters" section with dropdown menus for "Region" (set to "All") and "Cluster" (set to "All"), and a checkbox for "Show Anomalies Only" which is unchecked. The main content area has a title "Market Entry Advisor" with a blue circular icon. Below it is a "Business Use Case" section with the text "Planning expansion into new markets? Get data-driven recommendations on which markets to prioritize and a list of top target companies." There are two buttons: "Ask AI" (red) and "Use Filters" (grey). Underneath is a "Natural Language Query" section with the placeholder text "Example: 'Find fintech companies in Singapore with revenue over \$5M'". It includes a text input field with the query "I want to find manufacturing companies in Southeast Asia with more than 100 employees...", a "Search with AI" button (red), and a footer note "SDS Datathon 2026 - AI-Driven Company Intelligence Dashboard".

Figure 9: AI-Powered Market Entry Advisor

### 3.8.4 Competitive Benchmarking

## 4 Results

### 4.1 Exploratory Data Analysis

#### 4.1.1 Distribution Analysis

Table 4: Company Size Distribution

Employee Range	Count	Percentage
1-10	4,200	49%
11-50	2,100	25%
51-200	1,300	15%
201-1000	650	8%
1000+	309	4%

#### 4.1.2 Correlation Analysis

Key correlations observed:

- Revenue  $\leftrightarrow$  Employees:  $r = 0.65$
- Revenue  $\leftrightarrow$  Market Value:  $r = 0.72$
- Entity Score  $\leftrightarrow$  Revenue:  $r = 0.31$

### 4.2 Company Segments

K-Means clustering produced 5 distinct market segments (Table 5):

Table 5: Cluster Profiles

Tier	Name	Count	Med. Revenue	Med. Employees
1	Global HQ	507	\$45.2M	1,250
2	Subsidiary	2,012	\$3.1M	180
3	Subsidiary	1,834	\$850K	65
4	Local HQ	2,987	\$280K	22
5	Branch	1,219	\$12K	3

### 4.3 Lead Scoring Results

Table 6: Lead Tier Distribution

Tier	Score Range	Count	Percentage
Priority	75-100	3	0.04%
Hot	50-74	425	5.0%
Warm	30-49	2,891	33.8%
Cold	0-29	5,240	61.2%

### 4.4 Risk Detection Results

Table 7: Risk Detection Summary

Risk Type	Count
Shell Companies (Rule-based)	3,063
Statistical Anomalies (Isolation Forest)	428
High-Risk Entities ( $\geq 3$ flags)	244
Orphan Subsidiaries	89
Data Quality Issues	1,245

## 5 Discussion and Commercial Value

### 5.1 Strategic Insights

Our analysis identifies three high-value company segments:

1. **High-Revenue, High-Productivity Firms:** 159 companies with Revenue > \$100M
2. **Complex but Efficient:** 87 companies with Family Size > 100 and Rev/Emp > \$50K
3. **Outperforming SMBs:** 412 companies with Revenue < \$1M but Revenue\_vs\_Industry > 100%

### 5.2 Commercial Applications

The scalability of the proposed system supports diverse business use cases:

#### 5.2.1 Territory Planning for Sales Teams

**Scenario:** A B2B software company needs to assign sales territories across Asia.

**Solution:** By filtering for "Hot Leads" (Tier 2) and segmenting by Region, sales managers can

mathematically balance potential revenue across territories. The "Lead Score" prioritizes which 50 companies a rep should call first, replacing intuition with data-driven probability.

### 5.2.2 Pre-Acquisition Due Diligence

**Scenario:** Private Equity firms evaluating potential manufacturing targets.

**Solution:** The "Due Diligence Report" module instantly flags risks (e.g., Shell Company status, Orphan Subsidiary) and generates an AI-summarized financial health check. This reduces initial screening time from hours to seconds.

### 5.2.3 Competitive Benchmarking

**Scenario:** A mid-market CEO asks "How do we compare to peers?"

**Solution:** Using cluster baselines, the system calculates relative performance (e.g., "Revenue per Employee is 45% above the Tier 3 median"). This provides objective benchmarks for investor presentations and strategic planning.

## 5.3 Limitations

- Geographic bias toward Asian companies
- Missing financial data requires imputation (introduces uncertainty)
- Static clustering (could benefit from online learning)
- LLM responses depend on API availability

## 6 Conclusion

This project successfully developed a comprehensive company intelligence system that:

- Processed 8,559 company records with 72 attributes
- Engineered 15+ derived features
- Segmented companies into 5 interpretable tiers (Silhouette: 0.34)
- Implemented multi-factor lead scoring (0-100 scale)
- Detected 3,063 potential shell companies and 428 statistical anomalies
- Integrated LLM capabilities for natural language insights

The methodology demonstrates that systematic data processing, feature engineering, and machine learning can transform raw company data into actionable business intelligence, with clear commercial applications for sales, risk, and strategy functions.

## References

- Chen, H., Chiang, R. H., and Storey, V. C. (2012). Business intelligence and analytics: From big data to big impact. *MIS Quarterly*, 36(4):1165–1188.
- Gemini Team, Google (2023). Gemini: A family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*.
- Jain, A. K. (2010). Data clustering: 50 years beyond k-means. *Pattern Recognition Letters*, 31(8):651–666.
- Liu, F. T., Ting, K. M., and Zhou, Z.-H. (2008). Isolation forest. In *2008 Eighth IEEE International Conference on Data Mining*, pages 413–422. IEEE.

- MacQueen, J. et al. (1967). Some methods for classification and analysis of multivariate observations. In *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, pages 281–297. Oakland, CA, USA.
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., et al. (2011). Scikit-learn: Machine learning in python. *Journal of Machine Learning Research*, 12:2825–2830.
- Provost, F. and Fawcett, T. (2013). *Data Science for Business: What You Need to Know about Data Mining and Data-Analytic Thinking*. O'Reilly Media, Sebastopol, CA.
- Troyanskaya, O., Cantor, M., Sherlock, G., Brown, P., Hastie, T., Tibshirani, R., Botstein, D., and Altman, R. B. (2001). Missing value estimation methods for dna microarrays. *Bioinformatics*, 17(6):520–525.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. (2017). Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, pages 5998–6008.